

# Deep Learning Using TensorFlow



Dr. Ash Pahwa

---

Lesson 5: Linear Regression in TensorFlow

Lesson 5.2: Linear Regression in TensorFlow  
Multiple Variables



# Outline


---

- 2 Variable Regression
- Multi Variable Regression
  - Data Normalization
  - Dataset
  - Multi Variable Regression in R
  - Python: Un-Normalized Data
    - Linear Regression in Scikit-Learn
    - Linear Regression in TensorFlow (No Solution)
  - Python: Normalized Data
    - Linear Regression in Scikit-Learn
    - Linear Regression in TensorFlow

# 2 Variable Regression

## Generate the Dataset

```
#####  
# Generate data  
#  
number_of_points = 500  
x_point = []  
y_point = []  
m = 0.22  
c = 0.78  
for i in range(number_of_points):  
    x = np.random.normal (0.0, 0.5)  
    y = m*x + c + np.random.normal(0.0,0.1)  
    x_point.append([x])  
    y_point.append([y])  
  
plt.plot(x_point, y_point, 'o', label='Input Data')  
plt.legend()  
Out[22]: <matplotlib.legend.Legend at 0x1a8a68333c8>
```





# 2 Variable Regression

## Linear Regression in Scikit-Learn

```
#####  
# Linear Regression Using SKLearn function  
#  
  
linreg = linear_model.LinearRegression()  
  
linreg.fit(x_point, y_point)  
Out[27]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)  
  
print (linreg.intercept_)  
[ 0.78390055]  
  
print (linreg.coef_)  
[[ 0.20174445]]
```

### Scikit-Learn: Answer

```
print (linreg.intercept_)  
[ 0.78390055]  
  
print (linreg.coef_)  
[[ 0.20174445]]
```






### Regression Equation

$$y = 0.2017x + 0.7839$$

# 2 Variable Regression Result

```
print (linreg.intercept_)
[ 0.78390055]

print (linreg.coef_)
[[ 0.20174445]]
```



## TensorFlow Answer

```
Slope = [ 0.20152435]
Intercept = [ 0.78390092]
```

Regression Equation  
 $y = 0.2015x + 0.7839$



# Example#1

---

## Multi Variable Regression



# AutoMpg Dataset

	A	B	C	D	E	F	G	H	I	
1	carID	mpg	cylinders	displacement	horsepower	weight	acceleration	origin	age	
2	1	18	8	307	130	3504	12	1	23	
3	2	15	8	350	165	3693	11.5	1	23	
4	3	18	8	318	150	3436	11	1	23	
5	4	16	8	304	150	3433	12	1	23	
6	5	17	8	302	140	3449	10.5	1	23	
7	6	15	8	429	198	4341	10	1	23	
8	7	14	8	454	220	4354	9	1	23	
9	8	14	8	440	215	4312	8.5	1	23	
10	9	14	8	455	225	4425	10	1	23	
11	10	15	8	390	190	3850	8.5	1	23	
12	11	15	8	383	170	3563	10	1	23	
13	12	14	8	340	160	3609	8	1	23	
14	13	15	8	400	150	3761	9.5	1	23	
15	14	14	8	455	225	3086	10	1	23	
16	15	24	4	113	95	2372	15	3	23	



# Data Normalization: Standardization & Scaling

---





# Data Standardization and Scaling

---

- Standardization Data Variation
  - -3 to +3

$$z = \frac{\text{Data Value} - \text{Mean}}{\text{Standard Deviation}} = \frac{y - \mu}{\sigma}$$

- Scaling Data Variation
  - 0 to 1

$$y_i^j = \frac{x_i^j - \min_j}{\max_j - \min_j}$$

# Example

```
> normalize = function(x) {
+   return( (x-min(x))/(max(x)-min(x))) }
> data = c(124,3,311,341,298,136,23,75,5,51,822,364,663,444,999)
> (standard.data = scale(data))
      [,1]
[1,] -0.603086904
[2,] -0.994156118
[3,]  0.001292791
[4,]  0.098252100
[5,] -0.040722910
[6,] -0.564303180
[7,] -0.929516578
[8,] -0.761453776
[9,] -0.987692164
[10,] -0.839021223
[11,]  1.652833026
[12,]  0.172587571
[13,]  1.138948687
[14,]  0.431145729
[15,]  2.224892951
attr(,"scaled:center")
[1] 310.6
attr(,"scaled:scale")
[1] 309.4081
> (normalized.data = normalize(data))
[1] 0.121485944 0.000000000 0.309236948 0.339357430 0.296184739
0.133534137 0.020080321 0.072289157 0.002008032 0.048192771 0.822289157
[12] 0.362449799 0.662650602 0.442771084 1.000000000
>
```

	A	B	C	D	E	F	
1		#	Data	Standardization		Scaling	
2		1	124	-0.60		0.12	
3		2	3	-0.99		0.00	
4		3	311	0.00		0.31	
5		4	341	0.10		0.34	
6		5	298	-0.04		0.30	
7		6	136	-0.56		0.13	
8		7	23	-0.93		0.02	
9		8	75	-0.76		0.07	
10		9	5	-0.99		0.00	
11		10	51	-0.84		0.05	
12		11	822	1.65		0.82	
13		12	364	0.17		0.36	
14		13	663	1.14		0.66	
15		14	444	0.43		0.44	
16		15	999	2.22		1	
17							
18		Mean	310.60		Minimum	3	
19		StdDev	309.41		Maximum	999	
20							



# Multi Variable Regression in R

---

# Multi Variable Regression in R

```
> data <- read.csv("autompg.csv")
> #####
> y = data$mpg
> x1 = data$cylinders
> x2 = data$displacement
> x3 = data$horsepower
> x4 = data$weight
> x5 = data$acceleration
> x6 = data$origin
> x7 = data$age
> result <- lm(y ~ x3 + x4)
> summary(result)
```

Regression Equation: R: Un-Normalized Data

$$mpg = 45.6541 - 0.0472 * horsepower - 0.0058 * weight$$

Call:

lm(formula = y ~ x3 + x4)

Residuals:

Min	1Q	Median	3Q	Max
-11.1154	-2.7785	-0.3008	2.2864	16.6244

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	45.6540784	0.7943519	57.473	< 2e-16 ***
x3	-0.0471671	0.0111012	-4.249	2.69e-05 ***
x4	-0.0057880	0.0005031	-11.506	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.246 on 389 degrees of freedom

Multiple R-squared: 0.7051, Adjusted R-squared: 0.7036

F-statistic: 465.1 on 2 and 389 DF, p-value: < 2.2e-16

# Multi Variable Reg: Normalized Data

```
> #####  
> normalize = function(x) {  
+   return( (x-min(x))/(max(x)-min(x))) }  
> n.x3 = normalize(x3)  
> n.x4 = normalize(x4)  
> n.y = normalize(y)  
> result <- lm(n.y ~ n.x3 + n.x4)  
> summary(result)
```

Regression Equation: R: Normalized Data  
 $mpg = 0.6618 - 0.2284 * horsepower - 0.5372 * weight$

```
Call:  
lm(formula = n.y ~ n.x3 + n.x4)
```

```
Residuals:  
      Min       1Q   Median       3Q      Max   
-0.29251 -0.07312 -0.00791  0.06017  0.43749
```

```
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)      
(Intercept)  0.66180      0.01079   61.335 < 2e-16 ***  
n.x3         -0.22839      0.05375   -4.249 2.69e-05 ***  
n.x4         -0.53722      0.04669  -11.506 < 2e-16 ***  
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.1117 on 389 degrees of freedom  
Multiple R-squared:  0.7051, Adjusted R-squared:  0.7036  
F-statistic: 465.1 on 2 and 389 DF, p-value: < 2.2e-16
```



# Un-Normalized Data

# Dataset

```
#####
```

```
# 1. Load the Libraries
```

```
# 2. Read the Dataset
```

```
#
```

```
import tensorflow as tf
```

```
import numpy as np
```

```
import pandas as pd
```

```
data = pd.read_csv('autompg.csv')
```

```
data.head()
```

```
Out[12]:
```

	carID	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	1	18	8	307	130	3504	12.0	
1	2	15	8	350	165	3693	11.5	
2	3	18	8	318	150	3436	11.0	
3	4	16	8	304	150	3433	12.0	
4	5	17	8	302	140	3449	10.5	

	origin	age
0	1	23
1	1	23
2	1	23
3	1	23
4	1	23

$$mpg = \beta_0 + \beta_1 * horsepower + \beta_2 * weight$$

```
var_1 = data['horsepower']
```

```
var_1 = var_1.tolist()
```

```
var_2 = data['weight']
```

```
var_2 = var_2.tolist()
```

```
pred_vars_skl = data[['horsepower','weight']]
```

```
target = data['mpg']
```

```
target_skl = data['mpg']
```

```
target = target.tolist()
```



# Linear Regression in Scikit-Learn

---



# Implementation in Scikit-Learn

```
#####  
# 3. Compute the regression equation using Scikit-Learn  
#  
  
from sklearn import linear_model  
  
linreg = linear_model.LinearRegression()  
linreg.fit(pred_vars_skl, target)  
Out[26]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)  
  
print (linreg.intercept_)  
45.6540784147  
  
print (linreg.coef_)  
[-0.04716711 -0.00578798]
```

Regression Equation: R: Un-Normalized Data

$$mpg = 45.6541 - 0.0472 * horsepower - 0.0058 * weight$$

Regression Equation: Scikit-Learn: Un-Normalized Data

$$mpg = 45.6541 - 0.0472 * horsepower - 0.0058 * weight$$



# Linear Regression in TensorFlow

---



# Define the Hyper Parameters

---

```
#####  
# 4. Define the Hyper Parameters  
#  
learning_rate = 0.0000001  
epochs = 5000000
```



# Create Variables

## Build the Model

```
#####  
# 5. Create Variables X,Y,W,B  
# Build the model  
#  
  
x1 = tf.placeholder(dtype=np.float32)  
  
x2 = tf.placeholder(dtype=np.float32)  
  
y = tf.placeholder(dtype=np.float32)  
  
#####  
  
W1 = tf.Variable(np.random.randn(), dtype=np.float32, name="weight1")  
W2 = tf.Variable(np.random.randn(), dtype=np.float32, name="weight2")  
  
B = tf.Variable(np.random.randn(), dtype=np.float32, name="bias")  
  
#####  
computed_y = W1*x1 + W2*x2 + B
```



# Define the 'cost' and 'optimization' functions

---

```
#####  
# 6. Define the 'cost' and 'optimization' Functions  
# Initialize the variables  
#  
  
cost = tf.reduce_sum(tf.square(computed_y - y))  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)  
  
init = tf.global_variables_initializer()
```

# Run the Iteration Loop

```
with tf.Session() as sess:
    sess.run(init)
    for epoch in range(epochs):
        sess.run(optimizer, {x1:var_1, x2:var_2, y:target})

        if not epoch % 10000 :
            c = sess.run(cost, {x1:var_1, x2:var_2, y:target})
            w1 = sess.run(W1)
            w2= sess.run(W2)
            b = sess.run(B)
            #print(f'epoch={epoch: 4d} c={c: 4f} w1={w1: 4f} w2={w2: 4f} b={b: 4f}')
            print(epoch, c, w1, w2, b)

    weight1 = sess.run(W1)
    weight2 = sess.run(W2)
    bias = sess.run(B)
    print(f'Slope1={weight1: 4f} Slope2={weight2: 4f} Intercept={bias: 4f}')
```

0 1.06415e+15 -17.772 -531.53 0.928732  
10000 nan nan nan nan  
20000 nan nan nan nan  
30000 nan nan nan nan  
40000 nan nan nan nan



## Normalized Data

# Dataset

```
#####
```

```
# Linear Regression auto.csv
```

```
#
```

```
#####
```

```
# 1. Load the Libraries
```

```
# 2. Read the Dataset
```

```
#
```

```
import pandas as pd
```

```
import tensorflow as tf
```

```
import numpy as np
```

```
from sklearn import preprocessing
```

```
data = pd.read_csv('autompg.csv')
```

```
data.head()
```

```
Out[13]:
```

	carID	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	1	18	8	307	130	3504	12.0	
1	2	15	8	350	165	3693	11.5	
2	3	18	8	318	150	3436	11.0	
3	4	16	8	304	150	3433	12.0	
4	5	17	8	302	140	3449	10.5	

	origin	age
0	1	23
1	1	23
2	1	23
3	1	23
4	1	23

$$mpg = \beta_0 + \beta_1 * horsepower + \beta_2 * weight$$



# Normalize the Data

```
var_1 = preprocessing.minmax_scale(data['horsepower'])
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475:
DataConversionWarning: Data with input dtype int64 was converted to float64.
    warnings.warn(msg, DataConversionWarning)
```

```
var_1 = var_1.tolist()
```

$$mpg = \beta_0 + \beta_1 * horsepower + \beta_2 * weight$$

```
var_2 = preprocessing.minmax_scale(data['weight'])
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475:
DataConversionWarning: Data with input dtype int64 was converted to float64.
    warnings.warn(msg, DataConversionWarning)
```

```
var_2 = var_2.tolist()
```

```
pred_vars_sk1 = preprocessing.minmax_scale(data[['horsepower', 'weight']])
```

```
target_sk1 = preprocessing.minmax_scale(data['mpg'])
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475:
DataConversionWarning: Data with input dtype int64 was converted to float64.
    warnings.warn(msg, DataConversionWarning)
```

```
target_sk1 = target_sk1.tolist()
```

```
target = preprocessing.minmax_scale(data['mpg'])
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475:
DataConversionWarning: Data with input dtype int64 was converted to float64.
    warnings.warn(msg, DataConversionWarning)
```

```
target = target.tolist()
```



# Linear Regression in Scikit-Learn

---

# Implementation in Scikit-Learn

```
#####  
# 3. Compute the regression equation using Scikit-Learn  
#  
  
from sklearn import linear_model  
  
linreg = linear_model.LinearRegression()  
  
linreg.fit(pred_vars_skl, target_skl)  
Out[28]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)  
  
print (linreg.intercept_)  
0.661799255403  
  
print (linreg.coef_)  
[-0.22838813 -0.53721638]
```

Regression Equation: R: Normalized Data

$$mpg = 0.6618 - 0.2284 * horsepower - 0.5372 * weight$$

Regression Equation: Scikit-Learn: Normalized Data

$$mpg = 0.6618 - 0.2284 * horsepower - 0.5372 * weight$$



# Linear Regression in TensorFlow

---



# Define the Hyper Parameters

---

```
#####  
# 4. Define the Hyper Parameters  
#  
learning_rate = 0.0000001  
epochs = 5000000
```

# Create Variables

## Build the Model

```
#####  
# 5. Create Variables X,Y,W,B  
# Build the model  
#  
  
x1 = tf.placeholder(dtype=np.float32)  
x2 = tf.placeholder(dtype=np.float32)  
y = tf.placeholder(dtype=np.float32)  
  
#####  
W1 = tf.Variable([0], dtype=np.float32, name="weight1")  
W2 = tf.Variable([0], dtype=np.float32, name="weight2")  
B = tf.Variable([0], dtype=np.float32, name="bias")  
#####  
  
computed_y = W1*x1 + W2*x2 + B
```



# Define the 'cost' and 'optimization' functions

---

```
#####  
# 6. Define the 'cost' and 'optimization' Functions  
# Initialize the variables  
#  
  
cost = tf.reduce_sum(tf.square(computed_y - y))  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)  
  
init = tf.global_variables_initializer()
```

# Run the Iteration Loop

```
with tf.Session() as sess:
    sess.run(init)
    for epoch in range(epochs):
        sess.run(optimizer, {x1:var_1, x2:var_2, y:target})

        if not epoch % 100000 :
            c = sess.run(cost, {x1:var_1, x2:var_2, y:target})
            w1 = sess.run(W1)
            w2= sess.run(W2)
            b = sess.run(B)
            print(epoch, c, w1, w2, b)

weight1 = sess.run(W1)
weight2 = sess.run(W2)
bias = sess.run(B)
print(weight1, weight2, bias)
```

Regression Equation: TensorFlow: Normalized Data  
 $mpg = 0.6394 - 0.3084 * horsepower - 0.4183 * weight$

```
0 73.4795 [ 6.88996124e-06] [ 8.35379069e-06] [ 2.99000003e-05]
100000 10.8068 [-0.10149054] [-0.12765189] [ 0.44707638]
200000 6.75203 [-0.20843855] [-0.26567879] [ 0.54149044]
300000 5.48872 [-0.26535648] [-0.34467709] [ 0.59411216]
400000 5.09222 [-0.29446131] [-0.39071572] [ 0.62335241]
[-0.30839255] [-0.4183189] [ 0.63943666]
```





# Final Result

---

Regression Equation: R: Normalized Data

$$mpg = 0.6618 - 0.2284 * horsepower - 0.5372 * weight$$

Regression Equation: Scikit-Learn: Normalized Data

$$mpg = 0.6618 - 0.2284 * horsepower - 0.5372 * weight$$

Regression Equation: TensorFlow: Normalized Data

$$mpg = 0.6394 - 0.3084 * horsepower - 0.4183 * weight$$



# Summary

---

- 2 Variable Regression
- Multi Variable Regression
  - Data Normalization
  - Dataset
  - Multi Variable Regression in R
  - Python: Un-Normalized Data
    - Linear Regression in Scikit-Learn
    - Linear Regression in TensorFlow (No Solution)
  - Python: Normalized Data
    - Linear Regression in Scikit-Learn
    - Linear Regression in TensorFlow