

Deep Learning Using TensorFlow



Dr. Ash Pahwa

Lesson 6:
Gradient Descent & Backpropagation Algorithms
Lesson 6.1: Gradient Descent Algorithm



Outline

- What is a Gradient?
- What is Gradient Descent Algorithm?
- Gradient Descent Algorithm
 - Minimum of a 2-Variable Function
 - Minimum & Maximum of a 2-Variable Function
 - Minimum of 3 Variable Function
 - Solving Regression problem
 - Solving Regression problem – Iris Dataset

Lesson#4.1/Slide#11

Example#1: Neural Network: Data = XOR

5. Define Cost & Optimization Functions

- Cost Function
 - Cross Entropy Cost Function
- Gradient Descent
 - Feed “Learning Rate” as a parameter to the optimization function

```
#####  
# 5. Define Cost and Optimization Function  
# Cost = Cross Entropy Cost Function  
# Optimization = Gradient Descent  
#  
cost = tf.reduce_mean(-Y*tf.log(compOutput) - (1-Y)*tf.log(1-compOutput))  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

Lesson#4.1/Slide#29

Example#2: Neural Network: Data = Iris Define Cost and Optimization Function

- Cost Function
 - Cross Entropy Cost Function
- Gradient Descent
 - Feed "Learning Rate" as a parameter to the optimization function

```
cost      = tf.reduce_mean  
(tf.nn.softmax_cross_entropy_with_logits_v2(labels=outputs, logits=pred_tensor))  
  
updates = tf.train.GradientDescentOptimizer(0.01).minimize(cost)
```

Lesson#5.1/Slide#19

Example#1: Linear Regression: 2 Vars Build the TensorFlow Graph

```
graph = tf.Graph()
with graph.as_default():
    slope = tf.Variable(tf.random_uniform([1], -1.0, 1.0))
    intercept = tf.Variable(tf.zeros([1]))
    response = slope*x_point + intercept

#####
# Define Cost + Optimization Functions
cost = tf.reduce_mean(tf.square(response - y_point))
optimizer = tf.train.GradientDescentOptimizer(0.5).minimize(cost)
```

- $Residual = y_i - (mx_i + b)$
- $Residual^2 = (y_i - (mx_i + b))^2$
- $Residuals\ Sum\ of\ Squares = (RSS) = \sum_{i=1}^N (y_i - (mx_i + b))^2$

Lesson#5.2/Slide#31

Example#2: Linear Regression: Multi variables Define the 'cost' and 'optimization' functions

```
#####  
# 6. Define the 'cost' and 'optimization' Functions  
# Initialize the variables  
#  
  
cost = tf.reduce_sum(tf.square(computed_y - y))  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)  
  
init = tf.global_variables_initializer()
```



What is a Gradient?




Gradient Vector: 2 Variables

- Definition: 2 variables x, y
 - The gradient vector of a function $y=f(x)$
 - $\nabla y = \nabla f(x) = \frac{\partial f}{\partial x} = \frac{df}{dx}$
- Gradient property
 - Gradient vector gives the direction of fastest increase (or decrease) of function $f(x)$

Example: Gradient

- $y = x^2 + 6x$
 - $\frac{dy}{dx} = 2x + 6$
 - $x = -4$
 - $y = (-4)^2 + 6(-4) = 16 - 24 = -8$
 - $\frac{dy}{dx} = 2x + 6 = 2 * (-4) + 6 = -2$
 - Gradient at point $(-4, -8) = -2$






Gradient Vector: 3 Variables

- Definition: 3 variables x, y, z
 - The gradient vector of a function $z=f(x, y)$
 - $\nabla z = \nabla f(x, y) = \frac{\partial f}{\partial x} \mathbf{i} + \frac{\partial f}{\partial y} \mathbf{j}$
- Gradient vector has 2 properties
 - Gradient vector gives the direction of fastest increase (or decrease) of function $f(x, y)$
 - Gradient vector is orthogonal (perpendicular) to the level surface (contour)

Example#1: Gradient

- $z = f(x, y) = x^2y^3 - 4y$
- $\frac{\partial f}{\partial x} = 2xy^3$
- $\frac{\partial f}{\partial y} = 3x^2y^2 - 4$
- $\nabla f(x, y) = 2xy^3\mathbf{i} + (3x^2y^2 - 4)\mathbf{j}$
- $\nabla f(2, -1) = -4\mathbf{i} + 8\mathbf{j}$

$$f(x, y) = x^2y^3 - 4y$$





Example#1 : Gradient

- $z = f(x, y) = x^2y^3 - 4y$
- $\frac{\partial f}{\partial x} = 2xy^3$
- $\frac{\partial f}{\partial y} = 3x^2y^2 - 4$
- $\nabla f(x, y) = 2xy^3\mathbf{i} + (3x^2y^2 - 4)\mathbf{j}$
- $\nabla f(2, -1) = -4\mathbf{i} + 8\mathbf{j}$

- Gradient vector has 2 properties
 - Gradient vector gives the direction of fastest increase (or decrease) of function $f(x, y)$
 - Gradient vector is orthogonal (perpendicular) to the level surface (contour)

$$f(x, y) = x^2y^3 - 4y$$



R Code to Generate Plot

- $z = f(x, y) = x^2y^3 - 4y$
- $\frac{\partial f}{\partial x} = 2xy^3$
- $\frac{\partial f}{\partial y} = 3x^2y^2 - 4$
- $\nabla f(x, y) = 2xy^3\mathbf{i} + (3x^2y^2 - 4)\mathbf{j}$
- $\nabla f(2, -1) = -4\mathbf{i} + 8\mathbf{j}$



```
x = seq(-8,8,0.1)
y = seq(-8,8,0.1)
z1 = function (x,y) { (x^2)*(y^3) - 4*y }
z = outer(x,y,z1)

library(rgl)
persp3d(x,y,z,col="green")

contour(x,y,z)
text(0,-2,"Gradient(2,-1) = -4i + 8j",cex=2)
points(2,-1,pch=19)

# Since gradient = -4i + 8j
# End point of the gradient
# -4 + 2 = -2
# 8 - 1 = 7
arrows(2,-1, -2,7,lwd=3 )
abline(v=0)
abline(h=0)
```

$$f(x, y) = x^2y^3 - 4y$$



What is Gradient Descent Algorithm?



What is Gradient Descent Algorithm?

- Suppose a function $y = f(x)$ is given
 - Gradient Descent algorithm allows us to find the values of 'x' where the 'y' value becomes minimum or maximum values
- The Gradient Descent algorithm can be extended to any function with 2 or more variables ' $z = f(x, y)$ '

- Function $y=f(x)$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \parallel_{x^t}$

- Function $z=f(x,y)$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x and y
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial z}{\partial x} \parallel_{x^t, y^t}$
 - $y^{t+1} \leftarrow y^t - \eta \frac{\partial z}{\partial y} \parallel_{x^t, y^t}$



Finding Minimum of a Function Using Gradient Descent Algorithm


2 Variables (x, y) Function

Example 1

Find the value of 'x' where the value of 'y' is minimum

- $y = x^2$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x = 0$
 - $x = 0$
- To find the value of 'x'
 - Where the value of 'y' is minimum
 - Correct answer: $x = 0$



```
> x = seq(-5,5,0.1)
> y = x^2
> dy_dx = function (w1) { 2*w1 }
> plot(x,y,type='l',xlim=c(-10,10),ylim=c(-10,10))
> lines(x,dy_dx(x))
> lines(c(0,0),c(-10,10),lwd=3)
> lines(c(-10,10),c(0,0),lwd=3)
```



Example 1: R Code

- $y = x^2$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x = 0$
 - $x = 0$
- Gradient Descent Algorithm
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \big|_{x^t}$

```
> dy_dx = function (w1) { 2*w1 }
> xStart = 20
> learningRate = 0.01
> maxLimit = 300
> xStartHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   xStartHistory[i] = xStart
+   xStart = xStart - learningRate * dy_dx(xStart)
+ }
> plot(xStartHistory,type='l')
> lines(c(-10,10),c(0,0),lwd=3,col='red')
> lines(c(maxLimit,maxLimit),c(0,0),lwd=3,col='red')
> lines(c(0,maxLimit),c(0,0),lwd=3,col='red')
```



Example 1: Python Code

- $y = x^2$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x = 0$
 - $x = 0$
- Gradient Descent Algorithm
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \Big|_{x^t}$



```
import numpy as np
import matplotlib.pyplot as plt

def dy_dx(w1):
    return(2*w1)

xStart = 20
learning_rate = 0.01
maxLimit = 300
xStartHistory = np.zeros(maxLimit)

for i in range(maxLimit):
    xStartHistory[i] = xStart
    xStart = xStart - learning_rate * dy_dx(xStart)

plt.plot(xStartHistory)
```




Example 2

Find the value of 'x' where the value of 'y' is minimum

- $y = x^2 + 6x$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x + 6 = 0$
 - $x = -3$
- To find the value of 'x'
 - Where the value of 'y' is minimum
 - Correct answer: $x = -3$



```
> x = seq(-10,5,0.1)
> y = x^2 + 6*x
> dy_dx = function (w1) { 2*w1 + 6 }
> plot(x,y,type='l',xlim=c(-10,5),ylim=c(-10,5))
> lines(x,dy_dx(x))
> lines(c(0,0),c(-10,5),lwd=3)
> lines(c(-10,5),c(0,0),lwd=3)
> lines(c(-3,-3),c(-20,20),lwd=3,col='red')
```



Example 2

- $y = x^2 + 6x$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x + 6 = 0$
 - $x = -3$
- Gradient Descent Algorithm
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \Big|_{x^t}$

```
> dy_dx = function (w1) { 2*w1 + 6 }  
> xStart = 20  
> learningRate = 0.1  
> maxLimit = 100  
> xStartHistory = rep(0,maxLimit)  
> for ( i in 1:maxLimit )  
+ {  
+   xStartHistory[i] = xStart  
+   xStart = xStart - learningRate * dy_dx(xStart)  
+ }  
> plot(xStartHistory,type='l',ylim=c(-5,3))  
> lines(c(0,maxLimit),c(-3,-3),lwd=3,col='red')
```





Gradient Descent Algorithm

Parameters

- Parameters
 - Initial value of ' x '
 - Learning Rate
- If the choice of initial value of ' x ' and learning rate is different
 - The Gradient Descent algorithm may not converge



Finding Minimum + Maximum of a Function Using Gradient Descent Algorithm

2 Variables (x,y) Function


Example 3

Find the value of 'x' where the value of 'y' is

* Minimum

* Maximum

- $y = 3x^4 + 16x^3 + 18x^2$
- $\frac{dy}{dx} = 12x^3 + 48x^2 + 36x$
- To find the value of 'x'
 - Where the value of 'y' is local minimum
 - Correct answer: $x = 3, y = -27$




- To find the value of 'x'
 - Where the value of 'y' is local maximum
 - Correct answer: $x = 1, y = 27$

```
> x = seq(0,4,0.1)
> y = 3*x^4 - 16*x^3 + 18*x^2
> dy_dx = function (w1) { 12*w1^3 - 48*w1^2 + 36*w1 }
> plot(x,y,type='l',xlim=c(0,5),ylim=c(-50,50))
> #lines(x,dy_dx(x))
> lines(c(0,0),c(-40,40),lwd=3)
> lines(c(-10,10),c(0,0),lwd=3)
> lines(c(3,3),c(-40,40),lwd=3,col='red')
> lines(c(1,1),c(-40,40),lwd=3,col='red')
> text(1,42,"Max")
> text(3,42,"Min")
```



Example 3: Minimum

- $y = 3x^4 + 16x^3 + 18x^2$
- $\frac{dy}{dx} = 12x^3 + 48x^2 + 36x$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \Big|_{x^t}$

```
> dy_dx = function (w1) { 12*w1^3 - 48*w1^2 + 36*w1 }
> xStart = 2
> learningRate = 0.001
> maxLimit = 300
> xStartHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   xStartHistory[i] = xStart
+   xStart = xStart - learningRate * dy_dx(xStart) #
For minimum 'subtract'
+   ## xStart = xStart + learningRate * dy_dx(xStart)
# maximum 'add'
+ }
> plot(xStartHistory,type='l')
> lines(c(0,maxLimit),c(3,3),lwd=3,col='red')
```



- To find the value of 'x'
 - Where the value of 'y' is local minimum
 - Correct answer: $x = 3, y = -27$





Minimum and Maximum

Minimum

- Algorithm: **Minimum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \parallel_{x^t}$

Maximum


- Algorithm: **Maximum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t + \eta \frac{\partial y}{\partial x} \parallel_{x^t}$

- Change value = $\eta \frac{\partial y}{\partial x} \parallel_{x^t}$
- The only difference between the algorithm of **minimum** and **maximum** is that instead of **subtracting** we **add** the 'change' value to the 'x' value in the loop.


Example 3: Maximum

- $y = 3x^4 + 16x^3 + 18x^2$
- $\frac{dy}{dx} = 12x^3 + 48x^2 + 36x$
- Gradient Descent Algorithm: **Maximum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t + \eta \frac{\partial y}{\partial x} \parallel_{x^t}$

```
> dy_dx = function (w1) { 12*w1^3 - 48*w1^2 + 36*w1 }
> xStart = 0.01
> learningRate = 0.001
> maxLimit = 500
> xStartHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   xStartHistory[i] = xStart
+   ## xStart = xStart - learningRate * dy_dx(xStart)
+   # For minimum 'subtract'
+   xStart = xStart + learningRate * dy_dx(xStart)
+   # maximum 'add'
+ }
> plot(xStartHistory,type='l')
> lines(c(0,maxLimit),c(1,1),lwd=3,col='red')
```



- To find the value of 'x'
 - Where the value of 'y' is local maximum
 - Correct answer: $x = 1, y = 5$






Finding Minimum of a Function Using Gradient Descent Algorithm

3 Variables (x,y,z) Function

Example 4

Find the value of 'x' where the value of 'y' is
* Minimum

- $z = f(x, y) = x^2 + y^2 - 2x - 6y + 14$
- $\frac{\partial z}{\partial x} = 2x - 2$
- $\frac{\partial z}{\partial y} = 2y - 6$
- To find the value of 'x' and 'y'
 - Where the value of 'z' is local minimum
 - Correct answer: $x = 1, y = 3$
 - $\frac{\partial z}{\partial x} = 2x - 2 = 0; x = 1$
 - $\frac{\partial z}{\partial y} = 2y - 6 = 0; y = 3$


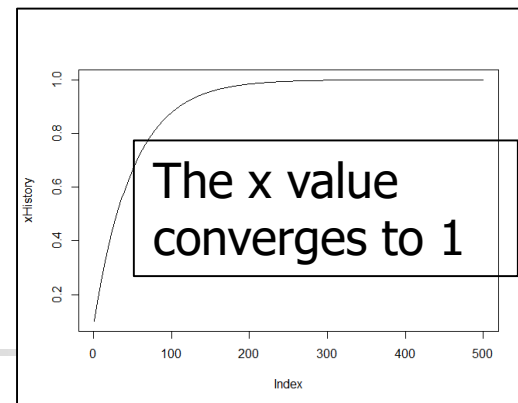


```
> x = seq(0,5,0.1)
> y = seq(0,10,0.1)
> z = function (x,y) { x^2 + y^2 - 2*x - 6*y + 14 }
> dz_dx = function (x1,y1) { 2*x1 - 2 }
> dz_dy = function (x1,y1) { 2*y1 - 6 }
> z<-outer(x,y,z)
> persp(x, y, z)
> contour(z)
```

Example 4: Minimum

- $z = f(x, y) = x^2 + y^2 - 2x - 6y + 14$
- $\frac{\partial z}{\partial x} = 2x - 2; \quad \frac{\partial z}{\partial y} = 2y - 6$
- To find the value of 'x' and 'y'
 - Where the value of 'z' is local minimum
 - Correct answer: $x = 1, y = 3$
 - $\frac{\partial z}{\partial x} = 2x - 2 = 0; x = 1$
 - $\frac{\partial z}{\partial y} = 2y - 6 = 0; y = 3$

```
> dz_dx = function (x1,y1) { 2*x1 - 2 }
> dz_dy = function (x1,y1) { 2*y1 - 6 }
> xStart = 0.1; yStart = 2
> learningRate = 0.01; maxLimit = 500
> xHistory = yHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit)
+ {
+   xHistory[i] = xStart
+   yHistory[i] = yStart
+   dW = dz_dx(xStart,yStart)
+   db = dz_dy(xStart,yStart)
+
+   xStart = xStart - learningRate * dW
+   yStart = yStart - learningRate * db
+ }
> plot(xHistory,type='l')
> plot(yHistory,type='l')
```



- Function $z=f(x,y)$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x and y
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial z}{\partial x} \big|_{x^t, y^t}$
 - $y^{t+1} \leftarrow y^t - \eta \frac{\partial z}{\partial y} \big|_{x^t, y^t}$

Copyright 2019 - Dr. Ash Pahwa

Solving Regression Problem Using Gradient Descent Algorithm




3 Variables (x, y, z) Function

Computing the Regression Line


Compute: Intercept and Slope

- Residual = Observed value – Computed Value
- Suppose regression equation is
 - $y = mx + b$
 - y is the explanatory variable
 - x is the predictor variable
 - m is the slope of the line
 - b is the intercept
- Residual = $y_i - (mx_i + b)$
- Residual² = $(y_i - (mx_i + b))^2$
- Residuals Sum of Squares = (RSS) = $\sum_{i=1}^N (y_i - (mx_i + b))^2$



Residual Sum of Squares

- *Residuals Sum of Squares* = (RSS) = $\sum_{i=1}^N (y_i - (mx_i + b))^2$
- To find the minimum point of this function,
 - we will take the partial derivative of RSS with respect to 'm' and 'b' and set that to zero.
- The RSS is a convex function and it has a minimum point



Partial Derivatives of the RSS w.r.t. Intercept and Slope

- *Residuals Sum of Squares* = (RSS) = $\sum_{i=1}^N (y_i - (mx_i + b))^2$
- To find the minimum point of this function,
 - we will take the partial derivative of RSS with respect to 'm' and 'b' and set that to zero.

$$\begin{aligned} \blacksquare \quad & RSS = \sum_{i=1}^N (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial b} = \sum_{i=1}^N \frac{\partial}{\partial b} (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial b} = -2 \sum_{i=1}^N (y_i - (mx_i + b)) \end{aligned}$$

$$\begin{aligned} \blacksquare \quad & RSS = \sum_{i=1}^N (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial m} = \sum_{i=1}^N \frac{\partial}{\partial m} (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial m} = -2 \sum_{i=1}^N (y_i - (mx_i + b))x_i \end{aligned}$$

$$\nabla RSS(b, m) = \begin{bmatrix} \frac{\partial RSS(m, b)}{\partial b} \\ \frac{\partial RSS(m, b)}{\partial m} \end{bmatrix} = \begin{bmatrix} -2 \sum_{i=1}^N (y_i - (mx_i + b)) \\ -2 \sum_{i=1}^N (y_i - (mx_i + b))x_i \end{bmatrix} = 0$$

Regression

Closed Form Solution

- To Compute 'm' and 'b'
 - SET GRADIENT = 0

- $$\nabla RSS(b, m) = \begin{bmatrix} \frac{\partial RSS(m, b)}{\partial b} \\ \frac{\partial RSS(m, b)}{\partial m} \end{bmatrix} = \begin{bmatrix} -2 \sum_{i=1}^N (y_i - (mx_i + b)) \\ -2 \sum_{i=1}^N (y_i - (mx_i + b)) x_i \end{bmatrix} = 0$$

- -----

- Top term

- $$b = \left(\frac{\sum y_i}{N} - m \frac{\sum x_i}{N} \right) = \mu_y - m\mu_x$$

- -----

- Bottom term

- $$m = \frac{\sum y_i x_i - \frac{\sum y_i \sum x_i}{N}}{\sum x_i^2 - \frac{\sum x_i \sum x_i}{N}} = r \frac{\sigma_y}{\sigma_x} = \text{Correlation} \frac{\text{Std Dev of } y}{\text{Std Dev of } x}$$

- -----

Closed Form Solution

$$m = \frac{\sum y_i x_i - \frac{\sum y_i \sum x_i}{N}}{\sum x_i^2 - \frac{\sum x_i \sum x_i}{N}}$$

$$b = \left(\frac{\sum y_i}{N} - m \frac{\sum x_i}{N} \right)$$

Regression Equation
 $y = 5x - 1$

	A	B	C	D	E	F	G
1							
2							
3		X	Y		X*Y		X^2
4		0	1		0		0
5		1	3		3		1
6		2	7		14		4
7		3	13		39		9
8		4	21		84		16
9							
10	SUM	10	45		140		30
11	AVERAGE	2	9		28		6
12	StdDev	1.58113883	8.124038405				
13	Correlation	0.97312368					
14							

C16 =E10-(B10*C10/5)							
	A	B	C	D	E	F	G
14							
15	Closed Form	Slope : Using SUM					
16		Numerator	50		(Sum of X*Y) - (1/N)*((Sum of X) * (Sum of Y))		
17		Denominator	10		(Sum of X^2) - (1/N)*((Sum of X * Sum of X))		
18		Slope	5				
19							
20		Intercept	-1		(Mean of Y) - slope * (Mean of X)		
21							

Regression

Using R 'lm' command

```
> x = c(0,1,2,3,4)
> y = c(1,3,7,13,21)
> plot(x,y)
> model = lm(y~x)
> summary(model)
```

Call:

```
lm(formula = y ~ x)
```

Residuals:

```
 1  2  3  4  5
2 -1 -2 -1  2
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.0000	1.6733	-0.598	0.59220
x	5.0000	0.6831	7.319	0.00527 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 2.16 on 3 degrees of freedom

Multiple R-squared: 0.947, Adjusted R-squared: 0.9293

F-statistic: 53.57 on 1 and 3 DF, p-value: 0.005268

```
> abline(model)
```

```
>
```



Regression Equation



$$y = 5x - 1$$

Regression

Gradient Descent Algorithm Approach

Regression Equation
 $y = 5x - 1$

```
> dRSS_dm = function (m,b) {-2*sum((y-m*x-b)*x) }
> dRSS_db = function (m,b) { -2*sum(y-m*x-b) }
> mStart = bStart = 0
> learningRate = 0.01;  maxLimit = 500
> mHistory = bHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   mHistory[i] = mStart
+   bHistory[i] = bStart
+
+   dW = dRSS_dm(mStart,bStart)
+   db = dRSS_db(mStart,bStart)
+
+   mStart = mStart - learningRate * dW
+   bStart = bStart - learningRate * db
+ }
> plot(mHistory,type='l')
> plot(bHistory,type='l')
```



Solving Regression Problem Using Gradient Descent Algorithm



Iris Dataset

Read the Iris Dataset

R Code


```
> data(iris)
> #####
> dim(iris)
[1] 150 5
> summary(iris)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.300	Min. :2.000	Min. :1.000	Min. :0.100	setosa :50
1st Qu.:5.100	1st Qu.:2.800	1st Qu.:1.600	1st Qu.:0.300	versicolor:50
Median :5.800	Median :3.000	Median :4.350	Median :1.300	virginica :50
Mean :5.843	Mean :3.057	Mean :3.758	Mean :1.199	
3rd Qu.:6.400	3rd Qu.:3.300	3rd Qu.:5.100	3rd Qu.:1.800	
Max. :7.900	Max. :4.400	Max. :6.900	Max. :2.500	

```
> head(iris)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
> x = iris$Petal.Length
> y = iris$Petal.Width
> plot(x,y)
```



Regression: R

Using R 'lm' command

```
> model = lm(y~x)
> summary(model)


Call:
lm(formula = y ~ x)

Residuals:
    Min       1Q   Median       3Q      Max
-0.56515 -0.12358 -0.01898  0.13288  0.64272

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.363076   0.039762  -9.131  4.7e-16 ***
x             0.415755   0.009582  43.387  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2065 on 148 degrees of freedom
Multiple R-squared:  0.9271, Adjusted R-squared:  0.9266
F-statistic: 1882 on 1 and 148 DF, p-value: < 2.2e-16

> abline(model)
```


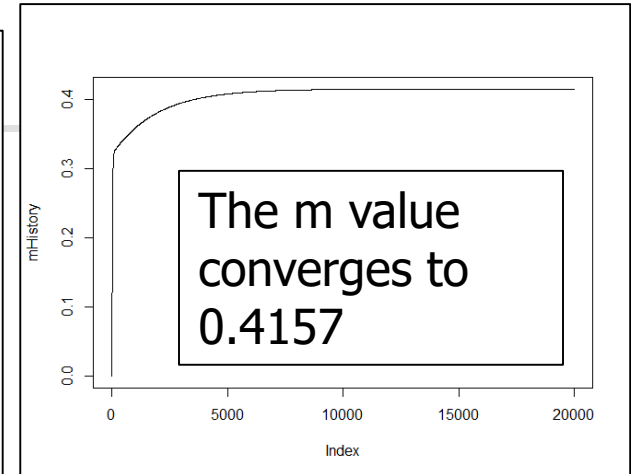


*Regression Equation: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*

Regression: R

Gradient Descent Algorithm Approach


```
> dRSS_dm = function (m,b) {-2*sum((y-m*x-b)*x) }
> dRSS_db = function (m,b) { -2*sum(y-m*x-b) }
> mStart = bStart = 0
> learningRate = 0.00001;    maxLimit = 20000
> mHistory = bHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   mHistory[i] = mStart
+   bHistory[i] = bStart
+
+   dW = dRSS_dm(mStart,bStart)
+   db = dRSS_db(mStart,bStart)
+
+   mStart = mStart - learningRate * dW
+   bStart = bStart - learningRate * db
+ }
> plot(mHistory,type='l')
> plot(bHistory,type='l')
> mHistory[maxLimit]
[1] 0.4157522
> bHistory[maxLimit]
[1] -0.3630608
```



*Regression Equation: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*
*Regression Eq: R Grad Desc: Petal.Width = 0.4157 * Petal.Length – 0.3631*

Read the Iris Dataset Python Code

```
#####  
# Load Libraries  
#  
from sklearn import linear_model  
from sklearn import datasets  
import matplotlib.pyplot as plt  
#####  
# 2. Read the Dataset  
#  
iris = datasets.load_iris()  
  
features = iris["data"]  
  
petalLength = features[:,2]  
petalLength[0:5]  
Out[13]: array([ 1.4,  1.4,  1.3,  1.5,  1.4])  
  
petalWidth = features[:,3]  
petalWidth[0:5]  
Out[15]: array([ 0.2,  0.2,  0.2,  0.2,  0.2])  
  
plt.plot(petalLength,petalWidth,'o')  
Out[16]: [<matplotlib.lines.Line2D at 0x16b604b01d0>]
```



Regression

Using Python 'Scikit-Learn' Library

```
#####  
# 3. Compute the regression equation using Scikit-Learn  
#  
petalLength = petalLength.reshape(-1,1)  
  
petalWidth = petalWidth.reshape(-1,1)  
  
linreg = linear_model.LinearRegression()  
  
linreg.fit(petalLength, petalWidth)  
Out[23]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1,  
normalize=False)  
  
print (linreg.intercept_)  
[-0.36651405]  
  
print (linreg.coef_)  
[[ 0.41641913]]
```

*Regression Equation: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Eq: R Grad Desc: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Equation: Scikit: Petal.Width = 0.4164 * Petal.Length – 0.3665*

Regression: Python

Gradient Descent Algorithm Approach

```
#####
```

```
# 1. Load the libraries
```

```
#
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn import datasets
```

```
#####
```

```
# 2. Read the Dataset
```

```
#
```

```
iris = datasets.load_iris()
```

```
features = iris["data"]
```

```
x = petalLength = features[:,2]
```

```
x[0:5]
```


```
Out[13]: array([ 1.4,  1.4,  1.3,  1.5,  1.4])
```

```
y = petalWidth = features[:,3]
```

```
y[0:5]
```

```
Out[15]: array([ 0.2,  0.2,  0.2,  0.2,  0.2])
```



```
plt.plot(x,y,'o')
```



Regression: Python

Gradient Descent Algorithm Approach

```
def dRSS_dm(m,b):  
    return (-2*sum((y-m*x-b)*x))  
  
def dRSS_db(m,b):  
    return (-2*sum((y-m*x-b)))  
  
mStart = 0  
bStart = 0  
learning_rate = 0.00001  
maxLimit = 20000  
mHistory = np.zeros(maxLimit)  
bHistory = np.zeros(maxLimit)  
  
for i in range(maxLimit):  
    mHistory[i] = mStart  
    bHistory[i] = bStart  
    #print(mHistory[i], bHistory[i])  
  
    dW = dRSS_dm(mStart,bStart)  
    db = dRSS_db(mStart,bStart)  
  
    mStart = mStart - learning_rate * dW  
    bStart = bStart - learning_rate * db  
  
print("mHistory=",mHistory[maxLimit-1])  
mHistory= 0.416415833891  
  
print("bHistory=",bHistory[maxLimit-1])  
bHistory= -0.366499084269
```





Final Result

*Regression Eq: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Eq: R Grad Desc: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Eq: Scikit: Petal.Width = 0.4164 * Petal.Length – 0.3665*

*Regression Eq: Python Grad Desc: Petal.Width = 0.4164 * Petal.Length – 0.3665*

Lesson#5.2/Slide#31

Example#2: Linear Regression: Multi variables Define the 'cost' and 'optimization' functions

```
#####  
# 6. Define the 'cost' and 'optimization' Functions  
# Initialize the variables  
#  
  
cost = tf.reduce_sum(tf.square(computed_y - y))  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)  
  
init = tf.global_variables_initializer()
```




Other Optimization Algorithms

- Stochastic Gradient Descent
- Momentum
- Nesterov Momentum
- AdaGrad
- RMSProp
- Adam: Adaptive Moments



Other Optimization Algorithms

Neural Network Optimization Algorithms

A comparison study based on TensorFlow

Vadim Smolyakov

<https://towardsdatascience.com/neural-network-optimization-algorithms-1a44c282f61dBy>

=====

Types of Optimization Algorithms used in Neural Networks and
Ways to Optimize Gradient Descent

Anish Singh Walia

<https://towardsdatascience.com/types-of-optimization-algorithms-used-in-neural-networks-and-ways-to-optimize-gradient-95ae5d39529f>



Summary

- What is a Gradient?
- What is Gradient Descent Algorithm?
- Gradient Descent Algorithm
 - Minimum of a 2-Variable Function
 - Minimum & Maximum of a 2-Variable Function
 - Minimum of 3 Variable Function
 - Solving Regression problem
 - Solving Regression problem – Iris Dataset
 - R + Python