

# Deep Learning Using TensorFlow



Dr. Ash Pahwa

---

Lesson 7:

Convolution Neural Network (CNN)

Lesson 7.2: MNIST Image Classification Application Using  
Linear Regression



# Outline

---

## ■ Step#1

1. Load Libraries
2. Download and Read MNIST data
3. One-hot Encoding
4. Function – Plotting images
5. Plot a Few images

## ■ Step#2

1. Placeholder variables
2. Variables
3. Model
4. Cost Function
5. Optimization Function
6. Performance Measures

## ■ Step#3

1. Create TensorFlow Session
2. Initialize Variables
3. Function: Optimization Iteration
4. Function: Display Performance
5. Function: Plot Model Weights

## ■ Step#4

1. Performance iteration count = 0
2. Performance iteration count = 1
3. Performance iteration count = 10
4. Performance iteration count = 1000
5. Print Confusion Matrix



# Step# 1

---

1. Load Libraries
2. Download and Read MNIST data
3. One-hot Encoding
4. Function – Plotting images
5. Plot a Few images



# 1.1 Load Libraries

---

```
#####  
# MNIST Image Classification Using Linear Regression  
#  
#####  
# 1.1 Load the libraries  
#  
  
import matplotlib.pyplot as plt  
import tensorflow as tf  
import numpy as np  
from sklearn.metrics import confusion_matrix  
  
print(tf.__version__)  
1.9.0
```


# 1.2 Download and Read MNIST data

## 12 MB of data

```
#####  
# 1.2 Download and read MNIST data  
#  
from tensorflow.examples.tutorials.mnist import input_data  
data = input_data.read_data_sets("MNIST_data/", one_hot = True)
```


```
print("The Training Set is:")  
The Training Set is:  
print(data.train.images.shape)  
(55000, 784)  
print(data.train.labels.shape)  
(55000, 10)  
print("The Test Set is:")  
The Test Set is:  
print(data.test.images.shape)  
(10000, 784)  
print(data.test.labels.shape)  
(10000, 10)  
print("The Validation Set is:")  
The Validation Set is:  
print(data.validation.images.shape)  
(5000, 784)  
print(data.validation.labels.shape)  
(5000, 10)
```

		Count	Image Shape	Labels
1	Training	55,000	784 (28x28 pixels image)	Count = 10: 0-9
2	Test	10,000	784 (28x28 pixels image)	Count = 10: 0-9
3	Validation	5,000	784 (28x28 pixels image)	Count = 10: 0-9
	Total	70,000		




# Image Dimensions

```
#####  
# the images are stored in one-dimensional arrays of this length.  
#  
  
img_size_flat = data.train.images[0].shape[0]  
img_size_flat  
Out[34]: 784  
  
# Tuple with height and width of images used to reshape arrays.  
img_shape = (28,28)  
  
# Number of classes, one class for each of 10 digits.  
num_classes = 10
```



# 1.3 One-hot Encoding

```
#####  
# 1.3 One hot encoding  
#  
data.test.labels[0:5,:]  
Out[42]:  
array([[ 0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.],  
       [ 0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.],  
       [ 0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.],  
       [ 1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.],  
       [ 0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.]])  
  
data.test.cls = np.array([label.argmax() for label in data.test.labels])  
  
data.test.cls[0:5]  
Out[44]: array([7, 2, 1, 0, 4], dtype=int64)
```



# 1.4 Function – Plotting Images

```
def plot_images(images, cls_true , cls_pred=None):
    assert len(images) == len(cls_true) == 9

    # Create figure with 3x3 subplots.
    fig, axes = plt.subplots(3,3)
    fig.subplots_adjust(hspace=0.3, wspace=0.3)

    for i, ax in enumerate(axes.flat):
        # Plot image.
        ax.imshow(images[i].reshape(img_shape), cmap='binary')

        # Show true and predicted classes
        if cls_pred is None:
            xlabel = 'True : {0}'.format(cls_true[i])
        else:
            xlabel = 'True : {0}, Pred : {1}'.format(cls_true[i], cls_pred[i])

        ax.set_xlabel(xlabel)


        # Remove ticks from the plot
        ax.set_xticks([])
        ax.set_yticks([])

    # Ensure the plot is shown correctly with multiple plots
    # in a single Notebook cell.
    plt.show()
```



# 1.5 Plot a Few Images

```
#####  
# 1.5 Plot a few images  
# Get the first images from the Test-set.  
#  
images = data.test.images[0:9]  
  
# Get the true classes for those images.  
cls_true = [np.argmax(oh) for oh in data.test.labels[0:9] ]  
  
# Plot the images and labels using our helper-function above.  
plot_images(images, cls_true)
```





## Step#2

---

1. Placeholder variables
2. Variables
3. Model
4. Cost Function
5. Optimization Function
6. Performance Measures



# 2.1 Placeholder Variables

Where the Input Data will be Assigned During Run Time

---

```
#####  
# 2.1 Placeholder variables  
#  
  
x = tf.placeholder( tf.float32, [None, img_size_flat])  
  
y_true = tf.placeholder( tf.float32, [None, num_classes])  
  
y_true_cls = tf.placeholder( tf.int64, [None])
```

`x = [?, 784]`

`y_true = [?, 10]`

`y_true_cls = [?]`

Contains a batch of images

Contains true one-hot encoded data

Contains the number

# 2.2 Variables

## Weights and Bias

```
#####  
# 2.2 Variables  
#  
weights = tf.Variable(tf.zeros([img_size_flat, num_classes]))  
bias = tf.Variable(tf.zeros([num_classes]))
```

```
weights = [784, 10]  
bias = [10]
```

Neural Network:

- Input Layer has 784 neurons
- Output Layer has 10 neurons

# Functions: 'softmax' and 'argmax'

$$\text{Softmax} = \frac{e^y}{\sum_1^n e^{y_i}}$$

	C	D	E	F
				$e^{y_i}$
			$e^{y_i}$	$\frac{e^{y_i}}{\sum_1^n e^{y_i}}$
		Logits Scores	EXP(n)	Softmax
		2	7.39	0.7
		1	2.72	0.2
		0.1	1.11	0.1
	SUM	3.1	11.21	

```
import numpy as np
a = np.array([(8,2,7),(3,4,5)])

print(a)
[[8 2 7]
 [3 4 5]]

#####
print(a.sum(axis=0))
[11  6 12]

print(a.sum(axis=1))
[17 12]

#####
print(a.argmax(axis=0))
[0 1 0]

print(a.argmax(axis=1))
[0 2]
```



## 2.3 Model

---

```
#####  
# 2.3 Model  
#  
logits = tf.matmul(x, weights) + bias  
y_pred = tf.nn.softmax(logits)  
y_pred_cls = tf.argmax(y_pred, axis=1)
```

$x = [?, 784]$

$weights = [784, 10]$   
 $bias = [10]$

$logits = [?, 784] * [784, 10] + [10] = [?, 10]$



## 2.4 Cost Function

---

```
# 2.4 Cost Function
#
cross_entropy = tf.nn.softmax_cross_entropy_with_logits_v2( logits= logits, labels = y_true)
cost = tf.reduce_mean(cross_entropy)
```



## 2.5 Optimization Function

---

```
#####  
# 2.5 Optimization Function  
#  
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.5).minimize(cost)
```





## 2.6 Performance Measures

---

```
# 2.6 Performance measures
#
correct_prediction = tf.equal( y_pred_cls , y_true_cls)
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```



## Step#3

---

1. Create TensorFlow Session
2. Initialize Variables
3. Function: Optimization Iteration
4. Function: Display Performance
5. Function: Plot Model Weights



# 3.1 Create TensorFlow Session

---

```
#####  
# 3.1 Create TensorFlow Session  
#  
  
session = tf.Session()
```



## 3.2 Initialize Variables

---

```
#####  
# 3.2 Initialize Variables  
#  
  
session.run(tf.global_variables_initializer())
```



## 3.3 Function: Optimization Iteration

```
#####  
# 3.3 Function to perform optimization iteration  
#  
  
batch_size = 100  
  
def optimize(num_iterations):  
    for i in range(num_iterations):  
        # Get a batch of training examples.  
        # x_batch now holds a batch of images and  
        # y_true_batch are the true labels for those images.  
        x_batch, y_true_batch = data.train.next_batch(batch_size= batch_size)  
  
        # Put the batch into a dict with the proper names  
        # for placeholder variables in the TensorFlow graph.  
        # Note that the placeholder for y_true_cls is not set  
        # because it is not used during training.  
        feed_dict_train = {x : x_batch,  
                           y_true : y_true_batch}  
  
        # Run the optimizer using this batch of training data.  
        # TensorFlow assigns the variables in feed_dict train  
        # to the placeholder variables and then runs the optimizer.  
        session.run(optimizer, feed_dict = feed_dict_train)
```



## 3.4-1 Display Performance

---

```
#####  
# 3.4 Optimization Iteration  
#  
feed_dict_test = {  
    x : data.test.images,  
    y_true : data.test.labels,  
    y_true_cls : [np.argmax(label) for label in data.test.labels]  
}  
  
def print_accuracy():  
    # Use TensorFlow to compute the accuracy.  
    acc = session.run(accuracy , feed_dict= feed_dict_test)  
  
    # Print the accuracy.  
    print('Accuracy on Test-st : {0:.1%}'.format(acc))
```

# 3.4-2 Display Performance

```
def print_confusion_matrix():
    # Get the true classifications for the Test-set.
    cls_true = [np.argmax(label) for label in data.test.labels]

    # Get the predicted classifications for the Test-set.
    cls_pred = session.run(y_pred_cls, feed_dict = feed_dict_test)

    # Get the confusion matrix using sklearn.
    cm = confusion_matrix(y_true = cls_true,
                          y_pred = cls_pred)

    # Print the confusion matrix as text.
    print(cm)

    # Plot the confusion matrix as an image.
    plt.imshow(cm, interpolation = 'nearest', cmap = plt.cm.Blues)

    # Make various adjustments to the plot.
    plt.tight_layout()
    plt.colorbar()
    tick_marks = np.arange(num_classes)
    plt.xticks(tick_marks, range(num_classes))
    plt.yticks(tick_marks, range(num_classes))
    plt.xlabel('Predicted')
    plt.ylabel('True')

    # Ensure the plot is shown correctly with multiple plots
    # in a single Notebook cell.
    plt.show()
```

# 3.4-3 Display Performance

```
def plot_example_errors():
    # Use TensorFlow to get a list of boolean values
    # whether each test-image has been correctly classified,
    # and a list for the predicted class of each image.
    correct, cls_pred = session.run([correct_prediction,
                                     y_pred_cls],
                                    feed_dict = feed_dict_test)

    # Negate the boolean array.
    incorrect = (correct == False)

    # Get the images from the Test-set that have been
    # incorrectly classified
    images = data.test.images[incorrect]

    # Get the predicted classes for those images
    cls_pred = cls_pred[incorrect]

    # Get the true classes for those images.
    cls_true = [np.argmax(label) for label in data.test.labels[incorrect]]

    # Plot the first 9 images.
    plot_images(images = images[0:9],
                cls_true = cls_true[0:9],
                cls_pred = cls_pred[0:9])
```



```

def plot_weights():
    # Get the values for the weights from the TensorFlow variable
    w = session.run(weights)

    # Get the lowest and highest values for the weights.
    # This is used to correct the colour intensity across
    # the images so they can be compared with each other.
    w_min = np.min(w)
    w_max = np.max(w)

    # Create figure with 3x4 sub-plots,
    # where the last 2 sub-plots are unused.
    fig, axes = plt.subplots(3,4)
    fig.subplots_adjust(hspace=0.3, wspace=0.3)

    for i, ax in enumerate(axes.flat):
        # Only use the weights for the first 10 sub-plots
        if i < 10:
            # Get the weights for the i'th digit and reshape it.
            # Note that w.shape == (img_size_flat, 10)
            image = w[:, i].reshape(img_shape)

            # Set the label for the sub-plot.
            ax.set_xlabel('Weights : {0}'.format(i))

            # Plot the image.
            ax.imshow(image, vmin=w_min, vmax= w_max,
                      cmap = 'seismic')

            # Remove ticks from each sub-plot
            ax.set_xticks([])
            ax.set_yticks([])

    # Ensure the plot is shown correctly with multiple plots
    # in a single Notebook cell.
    plt.show()

```

## 3.5 Function: Plot Model Weights



## Step#4

---

1. Performance iteration count = 0
2. Performance iteration count = 1
3. Performance iteration count = 10
4. Performance iteration count = 1000
5. Print Confusion Matrix

# 4.1 Performance:

## Iteration count = 0




---

```
# 4.1 Performance before any optimization
#
print_accuracy()
Accuracy on Test-st : 9.8%
```

# 4.1 Performance:

## Iteration count = 0

```
plot_example_errors()
```



# 4.2 Performance:


## Iteration count = 1

```
#####  
# 4.2 Performance Iteration#1  
#  
# Number of iteration means how many of batchs are iterated  
#  
optimize(num_iterations= 1)  
  
print_accuracy()  
Accuracy on Test-st : 50.9%
```

## 4.2 Performance:

### Iteration count = 1

```
plot_example_errors()
```



# 4.3 Performance:


## Iteration count = 10

```
#####  
# 4.3 Performance Iteration #10  
# We've already performed 1 iteration.  
#  
optimize( num_iterations=9 )  
  
print_accuracy()  
Accuracy on Test-st : 70.7%
```

## 4.3 Performance:

### Iteration count = 10

```
plot_example_errors()
```





# 4.4 Performance:


## Iteration count = 1000

```
#####  
# 4.4 Performance Iteration #1000  
# We've already performed 10 iterations.  
#  
optimize(num_iterations=990)  
  
print_accuracy()  
Accuracy on Test-st : 91.8%
```

# 4.4 Performance:


## Iteration count = 1000

```
plot_example_errors()
```



## 4.5 Print Confusion Matrix

```
#####  
# 4.5 Print confusion Matrix  
#  
print_confusion_matrix()  
[[ 964    0    2    3    0    3    5    1    2    0]  
 [   0 1104    2    2    1    2    4    2   18    0]  
 [   11   10  887   24   12    1   13   14   50   10]  
 [    3    0   13  936    1   17    1   10   20    9]  
 [    4    1    2    2  910    0   12    2    8   41]  
 [   10    3    3   57    9  746   19    7   31    7]  
 [   14    3    3    2   10   12  909    3    2    0]  
 [    3    8   18    9    5    1    0  944    3   37]  
 [    9    5    3   25    9   25   10   14  859   15]  
 [   13    6    1   11   27    6    0   22    7  916]]
```





# Summary

---

## ■ Step#1

1. Load Libraries
2. Download and Read MNIST data
3. One-hot Encoding
4. Function – Plotting images
5. Plot a Few images

## ■ Step#2

1. Placeholder variables
2. Variables
3. Model
4. Cost Function
5. Optimization Function
6. Performance Measures

## ■ Step#3

1. Create TensorFlow Session
2. Initialize Variables
3. Function: Optimization Iteration
4. Function: Display Performance
5. Function: Plot Model Weights

## ■ Step#4

1. Performance iteration count = 0
2. Performance iteration count = 1
3. Performance iteration count = 10
4. Performance iteration count = 1000
5. Print Confusion Matrix